

Unity ECC

Unified Memory Protection Against Bit and Chip Errors

Dongwhee Kim[†], Jaeyoon Lee[†], Wonyeong Jung[†],

Michael B. Sullivan^{††}, Jungrae Kim[†]

2023.11.15

[†]Sungkyunkwan University, ^{††}NVIDIA Corporation

Executive Summary

☑ Contributions

- Performance by 8.2% (Multi-core)
- DRAM energy consumption by 8.0%
- DRAM chip size by 6.9%

☑ Maintaining the same level of reliability

☑ Key idea

- Unity ECC can correct single-chip errors and double-bit errors



Contents

I. **Background**

II. **Motivation**

III. **Unity ECC**


IV. **Evaluation**

V. **Backup Slides**


Error Correction Codes (ECC)

✓ ECC can detect and correct errors

- ECC generates codeword with additional redundancy



Syndrome	Error
(001)	(0000001)
(010)	(0000010)
(100)	(0000100)
(101)	(0100000)
(110)	(1000000)
(011)	(0010000)
(111)	(0001000)

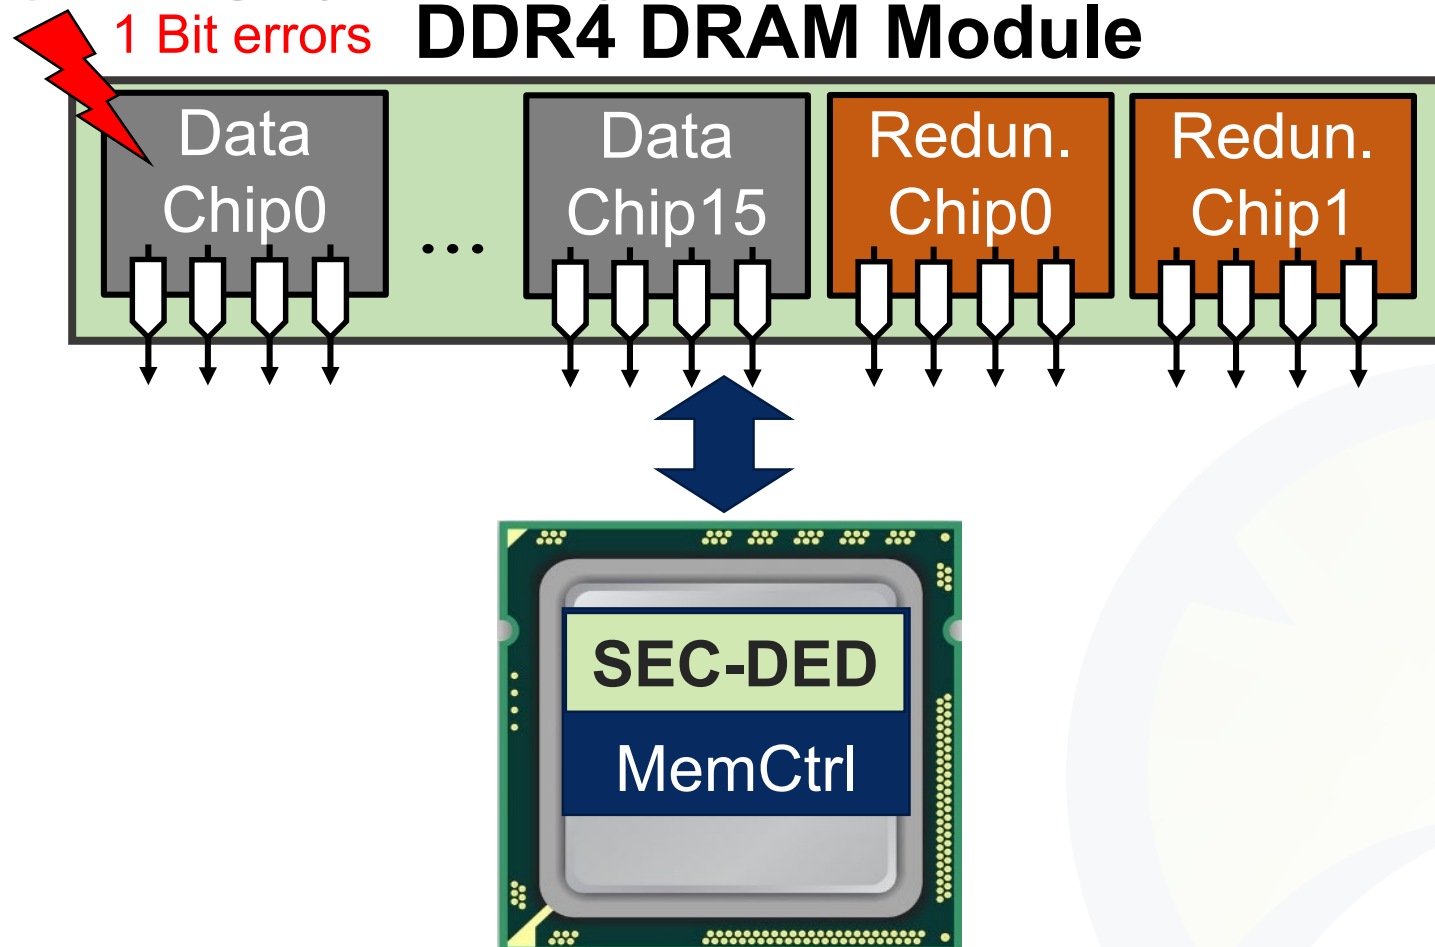


Single Error Correction-Double Error Detection

☑ **SEC-DED can correct single bit and detect double bit errors**

- Simple design, yet not very reliable

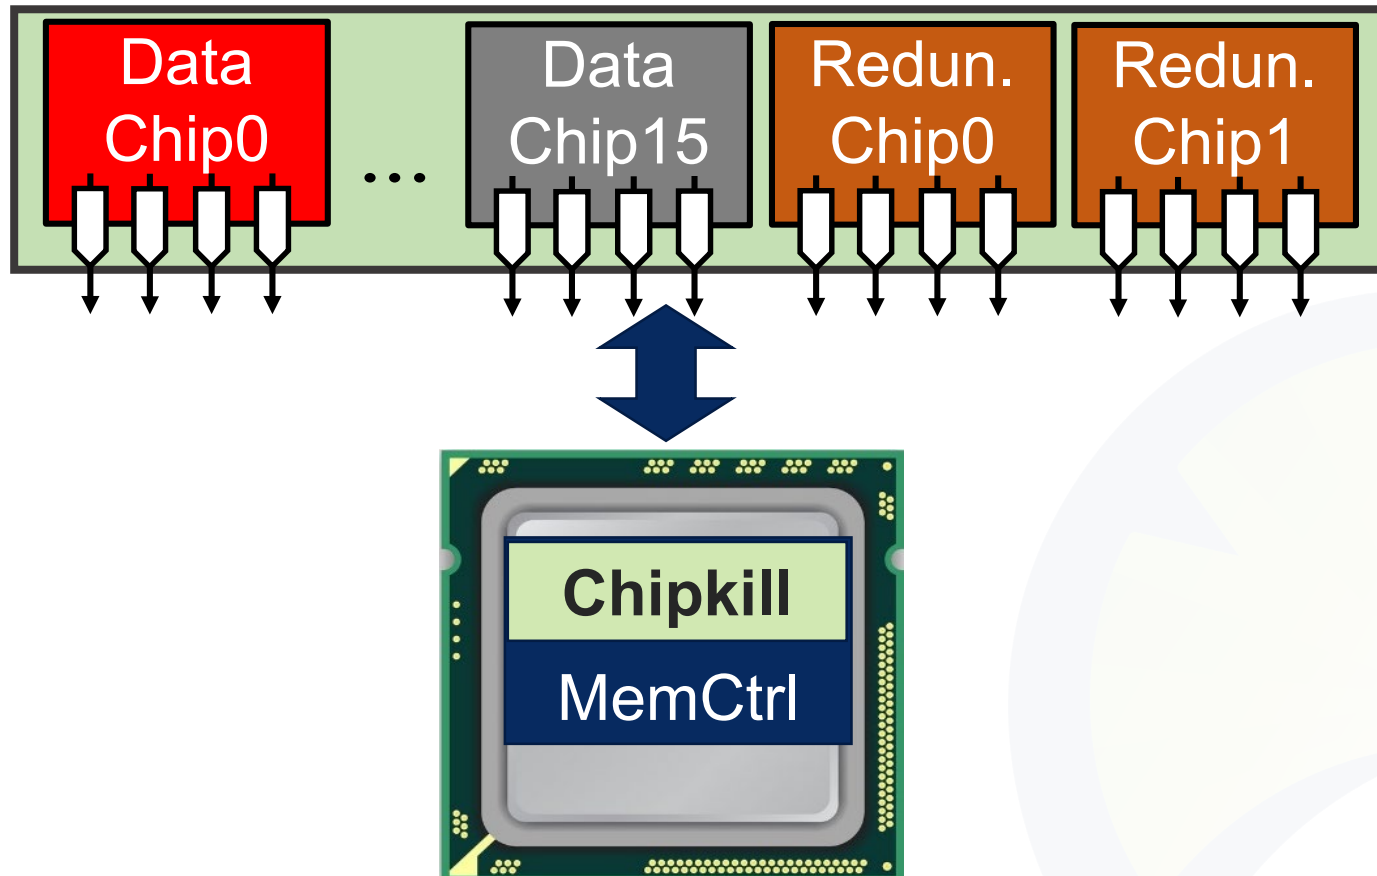
1 Bit errors **DDR4 DRAM Module**



☑ To correct more severe chip errors

- Susceptible to random bit errors

Chip errors **DDR4 DRAM Module**

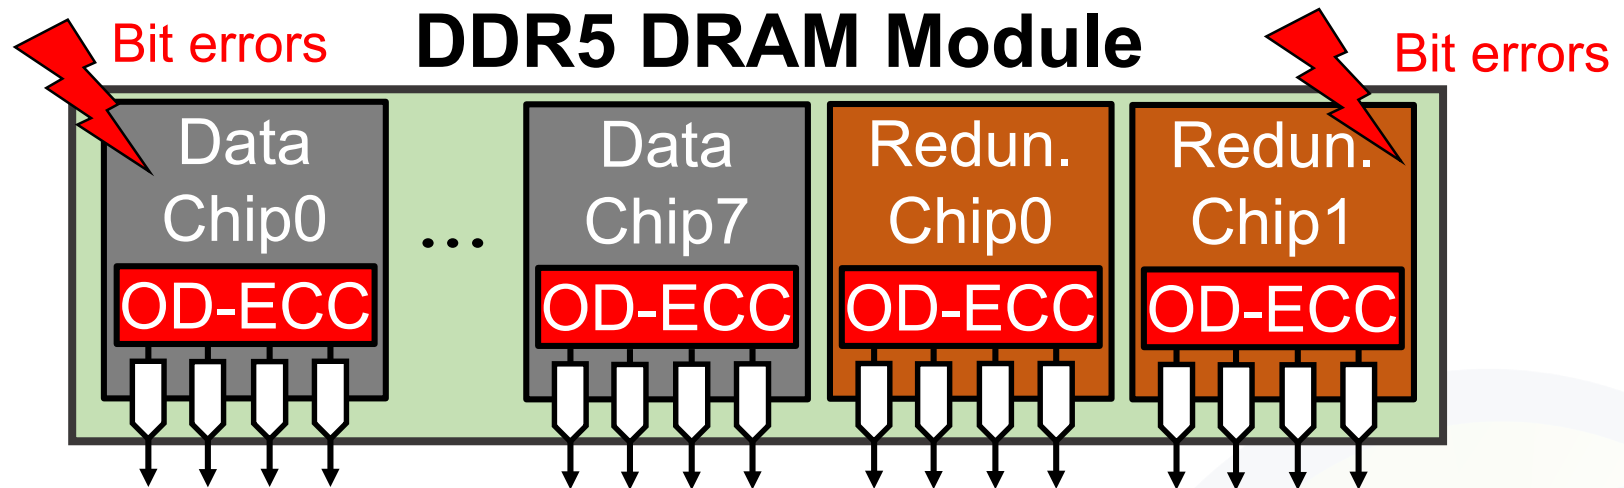


I. Background

DDR5 Configuration

☑ On-Die ECC (OD-ECC)

- OD-ECC can correct random bit errors



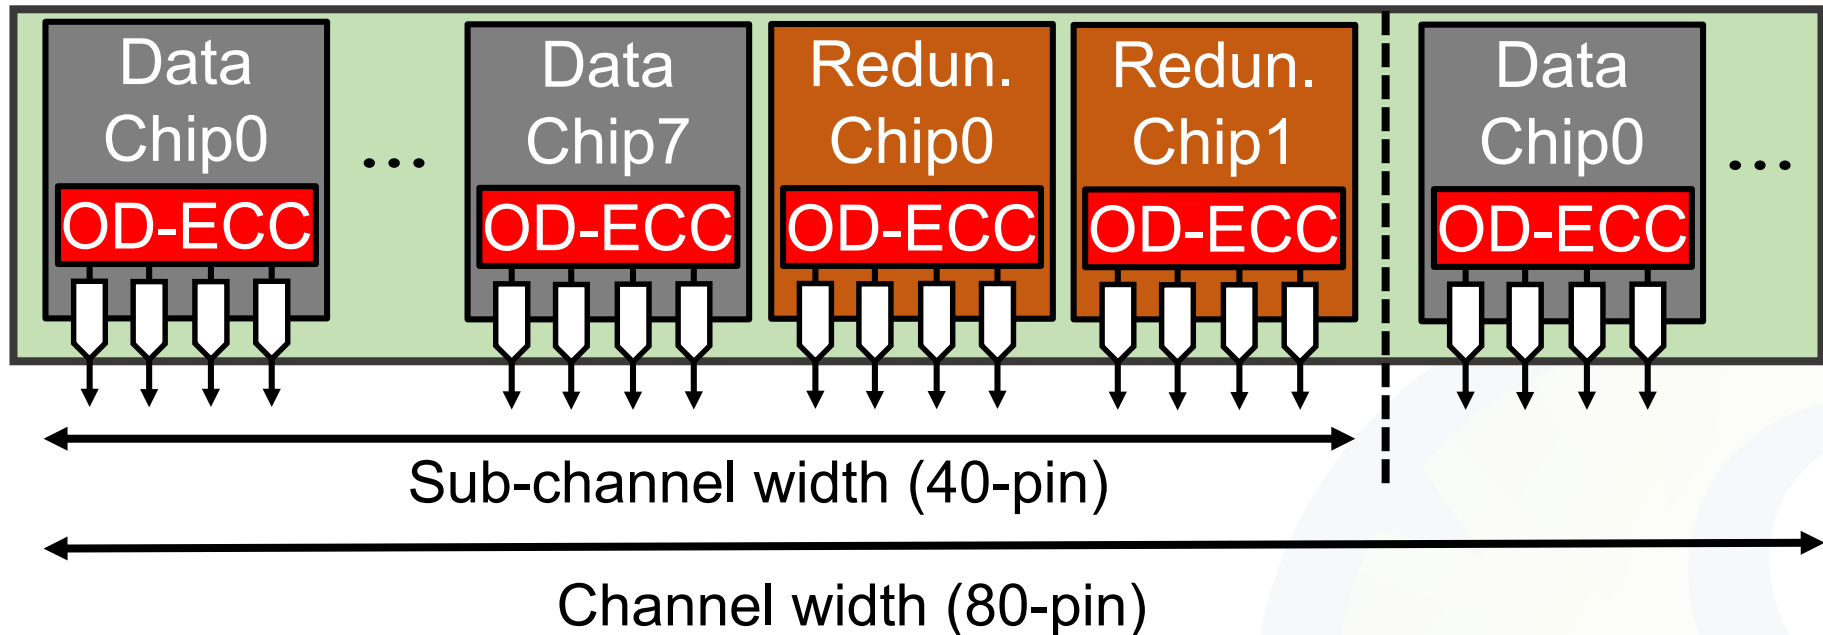
I. Background

DDR5 Configuration

☑ Sub-channel

- To align with the 64B cache granularity

DDR5 DRAM Module





Contents

I. Background

II. Motivation

III. Unity ECC

IV. Evaluation

V. Backup Slides

Cost of Two-level Protection

☑ High redundancy ratio (32.8%)

- OD-ECC (6.25%) X RL-ECC (25%)
 - RL-ECC (Rank-Level ECC): SEC-DED, Chipkill ...

☑ OD-ECC Overheads


- Performance
- DRAM energy
- DRAM chip area

II. Motivation

Shortened Code

✓ RL-ECC utilizes shortened code

- 8 bits per symbol (Reed-Solomon code)
- Unshortened code: **257** symbols (2,056 b)
- Shortened code: **10** symbols (80 b)



$$H = \begin{bmatrix} \begin{matrix} \text{8 columns} \\ \alpha^0 & \alpha^1 & \alpha^2 & \cdots & \alpha^{254} \\ \alpha^0 & \alpha^2 & \alpha^4 & \cdots & \alpha^{253} \end{matrix} & \begin{matrix} \text{2 columns} \\ \alpha^0 & 0 \\ 0 & \alpha^0 \end{matrix} \end{bmatrix}$$

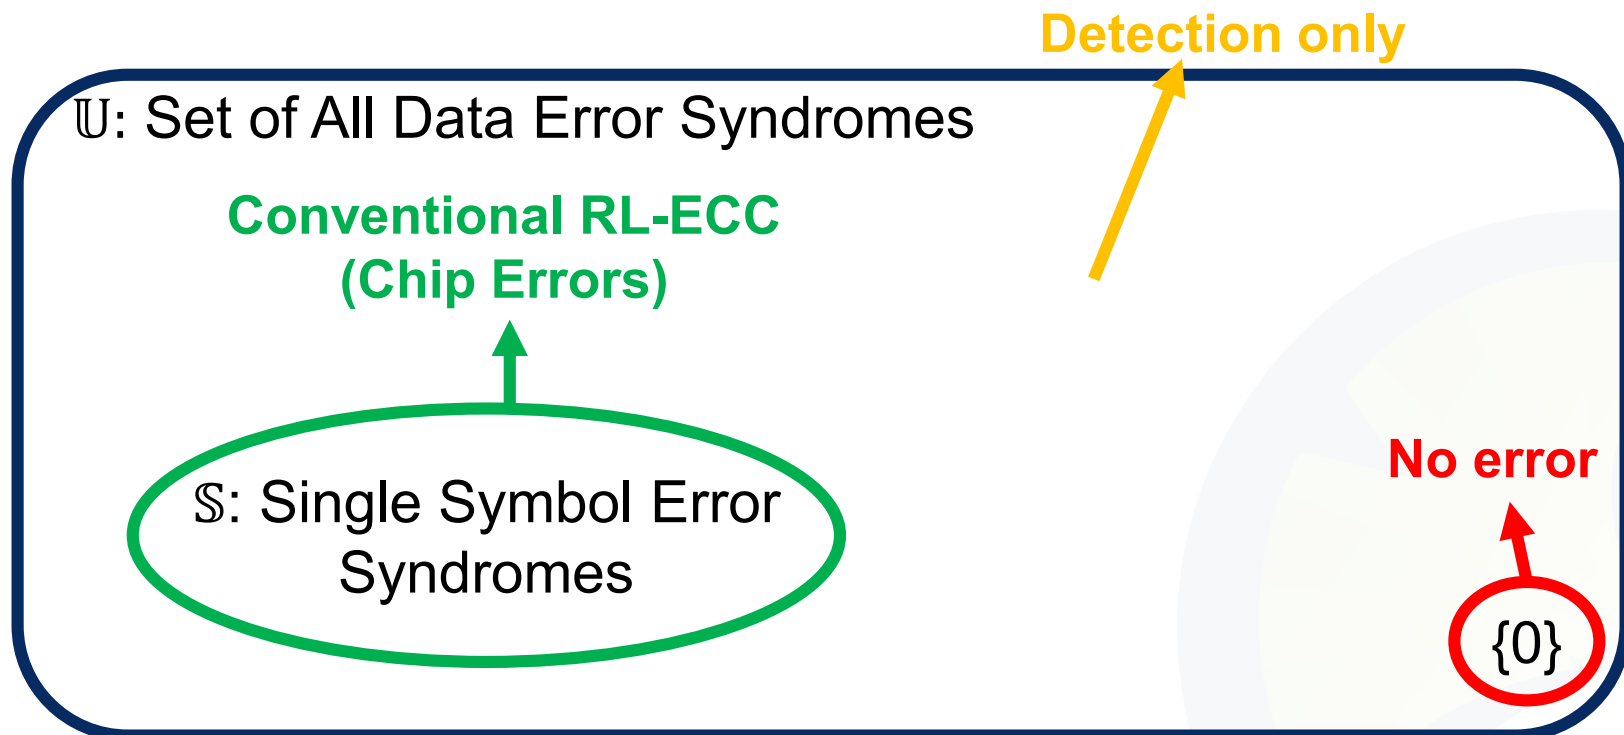
H-matrix
257 columns

$$H = \begin{bmatrix} \alpha^0 & \alpha^1 & \alpha^2 & \cdots & \alpha^7 & \alpha^0 & 0 \\ \alpha^0 & \alpha^2 & \alpha^4 & \cdots & \alpha^{14} & 0 & \alpha^0 \end{bmatrix}$$

H-matrix
10 columns

☑ Syndrome spaces of DDR5 RL-ECC

- Observation 1) **96.11%** of syndromes are used only for detection
- Observation 2) **Unused syndromes** can be employed for OD-ECC





Contents

I. Background

II. Motivation

III. **Unity ECC**

IV. Evaluation

V. Backup Slides

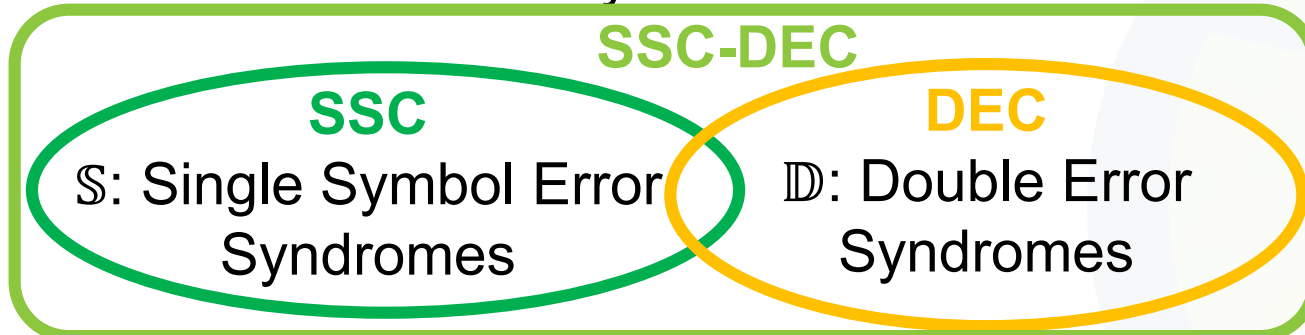
- ☑ **Unity ECC is a single-level RL-ECC (also Chipkill)**
- ☑ **Unity ECC can correct single chip errors and double bit errors**
 - SSC-DEC (Single Symbol Correction-Double Error Correction)
- ☑ **We propose flexible code construction algorithm**

Key idea: **Eliminate OD-ECC** by
remapping unused syndromes in RL-ECC

☑ H-matrix properties

- 1) All columns are **non-zero**
- 2) **DEC**: The sums (XOR operation) of any two columns are unique non-zero values
- 3) **SSC**: The sums (XOR operation) of all symbol-aligned columns are unique non-zero values
- 4) **SSC-DEC**: All sums from properties 2 and 3 should be unique


\mathbb{U} : Set of All Data Error Syndromes



No error
↑
{0}

✓ Based on Reed-Solomon (RS) code

- Systematic code
- Greedy search




257 columns

$$H = \begin{bmatrix} \alpha^0 & \alpha^1 & \alpha^2 & \cdots & \alpha^{254} & \begin{bmatrix} \alpha^0 & 0 \\ 0 & \alpha^0 \end{bmatrix} \\ \alpha^0 & \alpha^2 & \alpha^4 & \cdots & \alpha^{253} & \end{bmatrix}$$

Unshortened RS code H-matrix

- ✓ **Select 10 columns**
- ✓ Meet the H-matrix properties



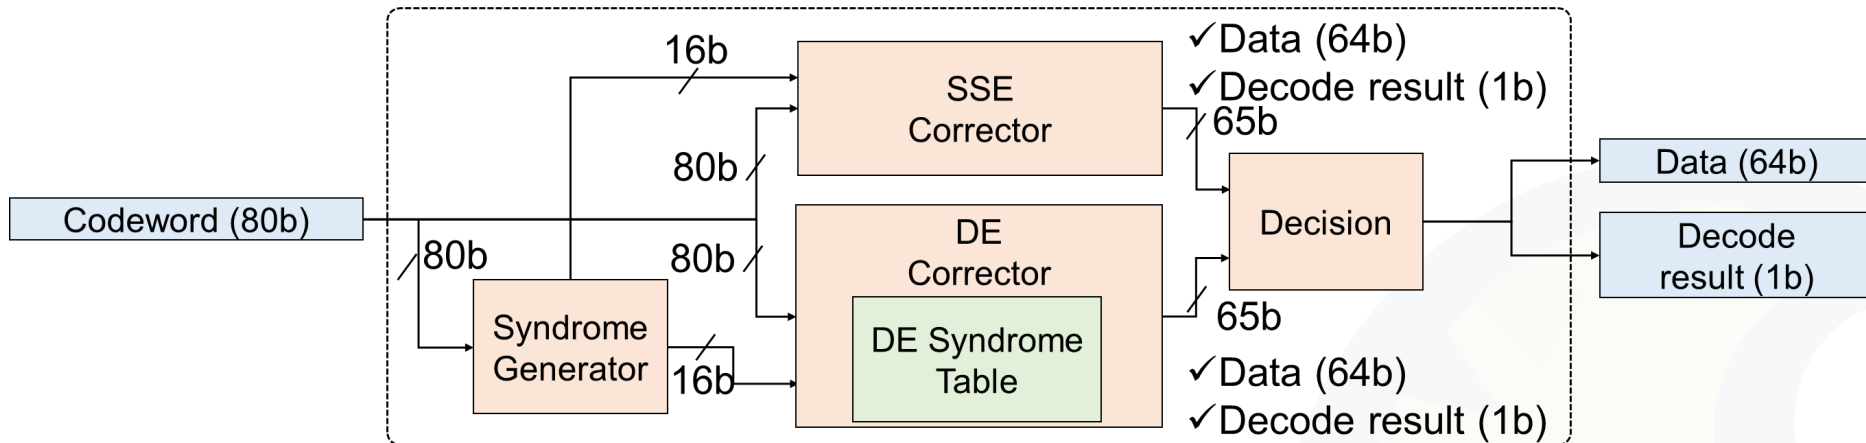
10 columns

$$H = \begin{bmatrix} \alpha^{25} & \alpha^{39} & \alpha^{63} & \alpha^{108} & \alpha^{141} & \alpha^{184} & \alpha^{215} & \alpha^{230} & \alpha^0 & 0 \\ \alpha^{50} & \alpha^{78} & \alpha^{126} & \alpha^{216} & \alpha^{27} & \alpha^{113} & \alpha^{175} & \alpha^{205} & 0 & \alpha^0 \end{bmatrix}$$

H-matrix example of (10, 8) **Unity ECC**

☑ Parallel decoder

- SSE Corrector (Single Symbol Error)
- DE Corrector (Double Error) [with DE Syndrome Table]





Contents

I. Background

II. Motivation

III. Unity ECC

IV. Evaluation

V. Backup Slides

IV. Evaluation

• Four Evaluation Metrics

☑ **Performance**

☑ **DRAM energy**

☑ **Reliability**

☑ **Hardware overheads**

- Synthesis (Synopsys Design Compiler, UMC 28nm)

☑ DRAM timing parameters

- Unity ECC eliminates OD-ECC (**lower DRAM timing parameters**)
- Baseline: OD-ECC + RL-ECC (Chipkill)
- Unity ECC: RL-ECC

Parameter	Baseline (OD-ECC + RL-ECC)	Unity ECC (RL-ECC)
Read latency (nCK)	40 (16.67ns)	36 (15ns)
tRCD (ns)	16.25	16.25
tRP (ns)	16.25	16.25
IDD4W (mA)	345	240

DRAM key timing parameters

☑ Simulation configuration

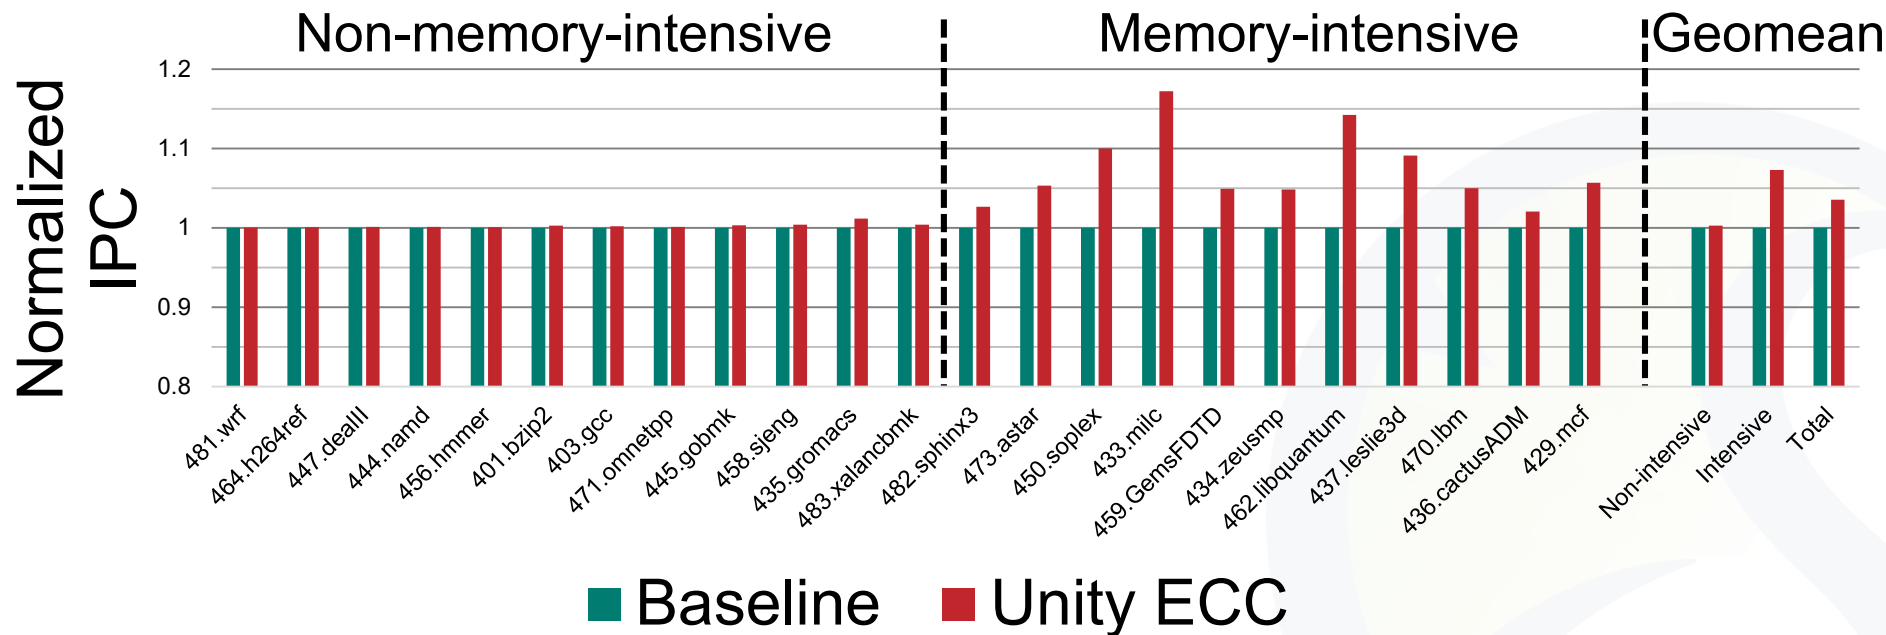
- Using Ramulator [1] and DRAMPower [2]
- Single-core/4-cores
- DRAM: DDR5-4800B, 16Gb, x4 chip, 32 banks
- Benchmarks: SPEC CPU2006 [4]

IV. Evaluation

Performance Results

☑ Single-core performance improvement

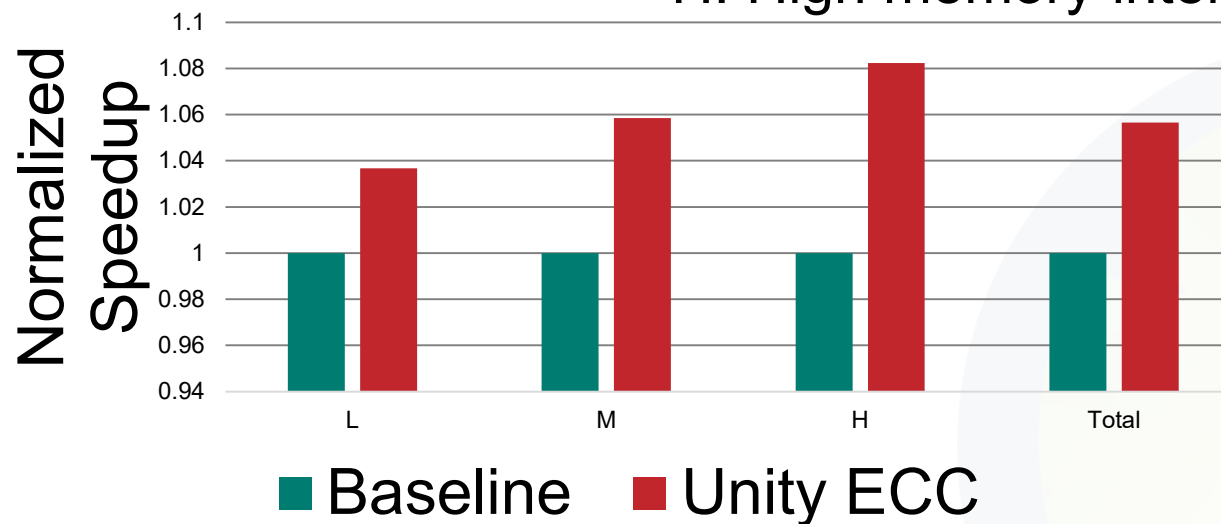
- 17.2% (max.)
- 7.27% (memory-intensive geomean)
- 3.56% (total geomean)



☑ Multi-core performance improvement

- 8.2% (max.)
- 5.7% (total geomean)

- ✓ L: Low memory intensity mix
- ✓ M: Medium memory intensity mix
- ✓ H: High memory intensity mix

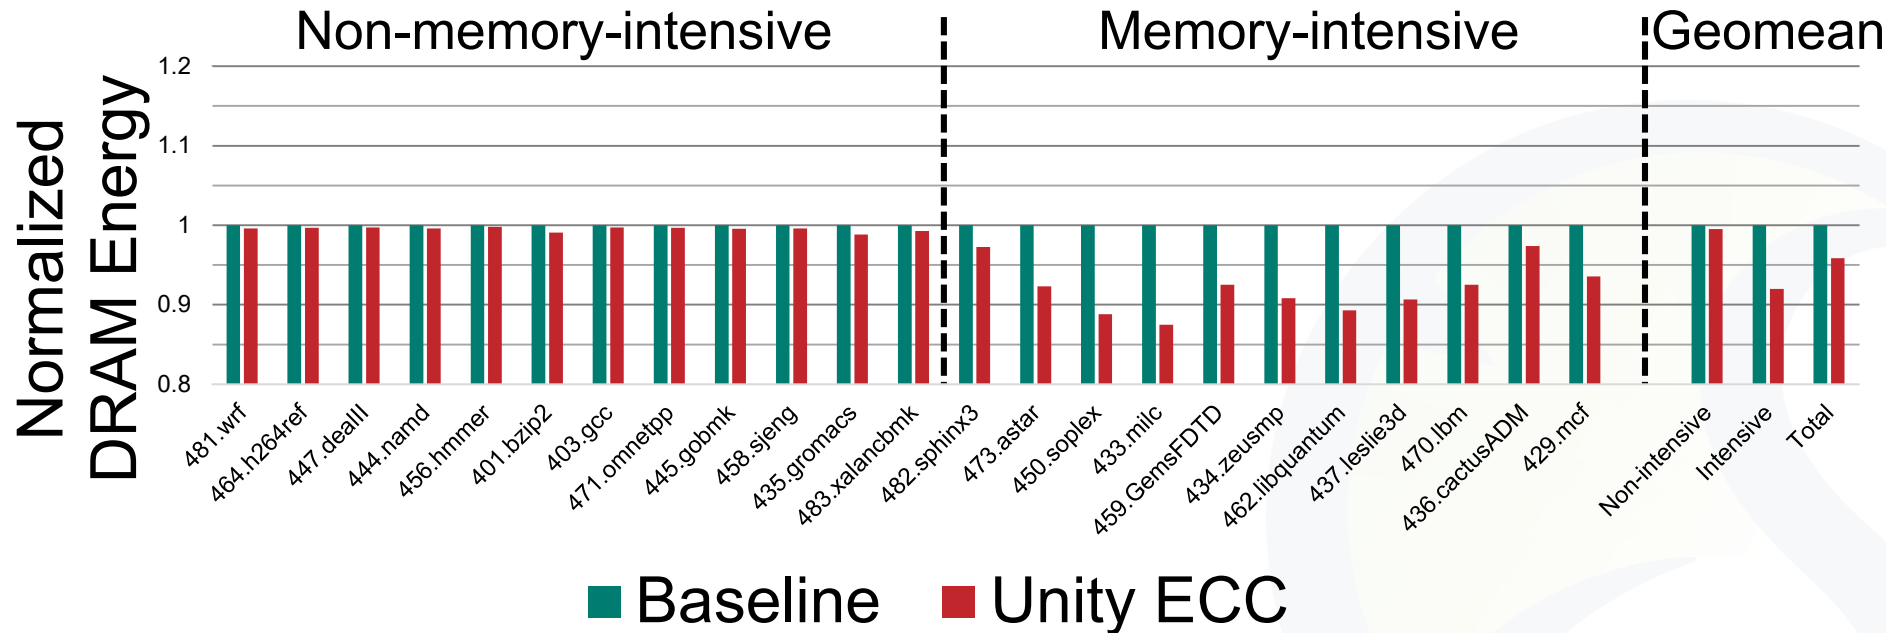


IV. Evaluation

DRAM Energy Results

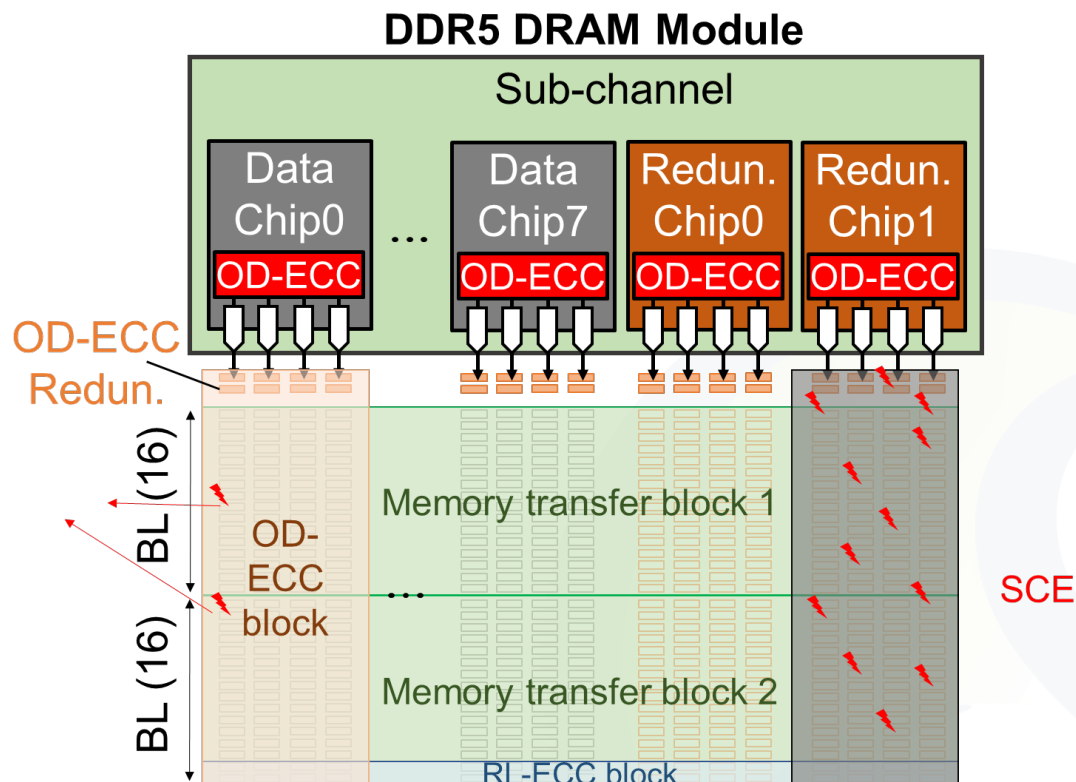
☑ Single-core DRAM energy reduction

- 12.53% (max.)
- 8.00% (memory-intensive geomean)
- 4.16% (total geomean)



Methodology

- Random injection of bit and chip errors
- The worse** of the two memory transfer blocks (128B) becomes the final output




- ✓ DBE: Double-Bit Error
- ✓ SCE: Single-Chip Error

IV. Evaluation

Reliability Results

☑ Bit errors

- Unity ECC \approx Baseline \gg OD-ECC $>$ Chipkill
 - Unity ECC even surpasses Baseline when $\text{BER} > 10^{-3}$



☑ Bit-error scenarios

- Unity ECC is robust to bit errors
 - Higher correction-probability of multi-bit errors
 - **Higher Correctable Error (CE) is better**

A comparison of reliability against bit-error scenarios (**CE cases**)

Error Scenario	OD-ECC	Chipkill	Baseline	Unity ECC
SBE (%)	100			
DBE (%)	0.14	100	100	100
DBE + DBE (%)	0.01	12.30	8.89	98.68

- ✓ SBE: Single-Bit Error
- ✓ DBE: Double-Bit Error

IV. Evaluation

Reliability Results

☑ Bit-error scenarios

- Unity ECC is robust to bit errors
 - Higher detection-probability of multi-bit errors
 - **Lower Silent Data Corruption (SDC) is better**

A comparison of reliability against bit-error scenarios (**SDC cases**)

Error Scenario	OD-ECC	Chipkill	Baseline	Unity ECC
DBE (%)	99.86	0	0	0
DBE + DBE (%)	99.99	1.03	1.75	0.09

- ✓ SBE: Single-Bit Error
- ✓ DBE: Double-Bit Error

☑ Chip errors

- Unity ECC \approx Baseline \approx Chipkill \gg OD-ECC
 - Unity ECC can correct all 1 chip errors
 - Unity ECC can detect most 2 chip errors (99.99999996%)

A comparison of reliability against chip-error scenarios (**CE cases**)

Error Scenario	OD-ECC	Chipkill	Baseline	Unity ECC
SCE (%)	0	100	100	100
SCE + SCE (%)	0	0	0	0

A comparison of reliability against chip-error scenarios (**SDC cases**)

SCE (%)	100	0	0	0
SCE + SCE (%)	100	0	0	0.00000004

✓ SCE: Single-Chip Error

☑ Chip errors

- Unity ECC \approx Baseline \approx Chipkill \gg OD-ECC
 - Unity ECC can correct all 1 chip errors

Unity ECC can detect most 2 chip errors (99.9999996%)

More results can be found in our paper

Scenario				
SCE (%)	0	100	100	100
SCE + SCE (%)	0	0	0	0

A comparison of reliability against chip-error scenarios (**SDC cases**)

SCE (%)	100	0	0	0
SCE + SCE (%)	100	0	0	0.0000004

✓ SCE: Single-Chip Error

☑ Unity ECC vs Baseline RL-ECC (Chipkill) – Latency

- Read latency slightly increases in rare error cases only
 - Doesn't affect normal read latency

	Encoder		Decoder	
	Baseline RL-ECC	Unity ECC	Baseline RL-ECC	Unity ECC
Encoding latency (ns)	0.25	0.25		
Decoding latency (ns) - detection			0.25	0.25
Decoding latency (ns) - overall			0.31	0.81

VI. Evaluation

Hardware Overheads Results

☑ Unity ECC vs Baseline RL-ECC (Chipkill) – Area and Power

- Decoder area slightly increases **0.009%** of the entire processor
- Power savings from DRAM can easily offset the 25.4mW
- Eliminating OD-ECC can reduce the DRAM chip size by **6.9%** [5]

	Encoder		Decoder	
	Baseline RL-ECC	Unity ECC	Baseline RL-ECC	Unity ECC
Area (μm^2)	113.06	264.26	1184.74	10135.78
Total power (mW)	0.34	1.02	3.75	29.13

Summary

☑ Observations

- OD-ECC overheads & Shortened code in RL-ECC

☑ Idea: Unity ECC

- Eliminate OD-ECC while maintaining reliability levels by repurposing unused correction syndromes

☑ Evaluation

- **Pros:** Performance (8.2%), DRAM energy (8.0%) and DRAM chip size (6.9%)
- **Cons:** Reliability and Hardware overheads (0.009%)

Reliability Evaluation Open Source

☑ ECC-ExerSim [3]

- Github search 'xyz123479/ECC-exercise'
- CERN-OHL-S-2.0 license

ECC-ExerSim



Thank you

Q&A





Contents

I. Background

II. Motivation

III. Unity ECC

IV. Evaluation

V. Backup Slides

☑ Comparison to Related Works

- Unity ECC has **high error correction capability**

- No additional redundancy (vs Chipkill)**

- ✓ n: codeword length
- ✓ k: data length

Class	Error Correction			ECC word configs		
	SE	DE	SSE	n	k	Redundancy
SEC	O	X	X	136	128	6.25%
DEC	O	O	X	144	128	12.5%
	O	O	X	78	64	21.875%
SSC (Chipkill)	O	X	O	80	64	25%
DEC-SbEC	O	O	O	88	64	37.5%
Unity ECC	O	O	O	80	64	25%

- ✓ SEC (Single Error Correcting)
- ✓ DEC (Double Error Correcting)
- ✓ SSC (Single Symbol Correcting)
- ✓ DEC-SbEC (Double Bit Error Correcting – Single b-bit Byte Error Correcting)

DRAM timing parameters

☑ DRAM timing parameters

- Due to OD-ECC

Parameter	Baseline (OD-ECC + Chipkill)	Unity ECC
Read latency (nCK)	40 (16.67ns)	36 (15ns)
tRCD (ns)	16.25	
tRP (ns)	16.25	
tCCD_L_WR (nCK)	48	24
Write latency (nCK)	38	34
tCCD_S_WTR (nCK)	52	48
tCCD_L_WTR (nCK)	70	66
IDD0 (mA)	103	
IDD4W (mA)	345	240

Simulation Configuration

☑ Simulation configuration

- Single-core
- Multi-core (4 core)

Processor	1 or 4 core(s), 4GHz, 4-wide issue, 8 MSHRs per core 128-entry instruction window
LLC	64B cacheline, 8-way associative, 47 CPU-cycle latency, 8MB total capacity (2MB/core with 4 cores)
Memory Controller	FR-FCFS-Cap scheduling, timer-based row open policy, 64-entry read/write request queue
DRAM	1 channel, 1 rank, DDR5 , 4800Mbps, 16Gb ×4 chip, 8 bank groups, 4 banks per bank group
Benchmarks	23 benchmarks from SPEC CPU2006

☑ Synthesis setup

- Compiler: Synopsys Design Compiler
- Logic libraries: UMC 28nm SVT/LVT cells (Choose the worst one)
- Clock uncertainty 60% (40% margin)
- Activity factor switching: 10%
- Corner: ss
- PVT (28nm, 0.9V, 125'C)

References

- [1] Yoongu Kim, Weikun Yang, and Onur Mutlu. 2015. Ramulator: A fast and extensible DRAM simulator. IEEE Computer architecture letters 15, 1 (2015), 45–49.
- [2] Karthik Chandrasekar, ChristianWeis, Yonghui Li, Benny Akesson, NorbertWehn, and Kees Goossens. 2012. DRAMPower: Open-source DRAM power energy estimation tool. URL: <http://www.drampower.info> 22 (2012).
- [3] <https://github.com/xyz123479/ECC-exercise>
- [4] SPEC CPU2006. 2006. Standard performance evaluation corporation

References

[5] Sanguhn Cha, O Seongil, Hyunsung Shin, Sangjoon Hwang, Kwangil Park, Seong Jin Jang, Joo Sun Choi, Gyo Young Jin, Young Hoon Son, Hyunyon Cho, et al. 2017. Defect analysis and cost-effective resilience architecture for future DRAM devices. In 2017 IEEE International Symposium on High Performance Computer Architecture (HPCA). IEEE, 61–72